

Pavel Donin

AI Engineer — Agent Orchestration, LLM Tooling

Agent systems in production for a studio team and client users · 10+ years in software, 5 of them leading teams

Contact

doninpr@gmail.com · linkedin.com/in/doninpr · t.me/doninpr · github.com/prineycom

Novi Sad, Serbia (CET, UTC+1) · remote · open to full-time roles worldwide

Professional Summary

AI engineer who builds agent systems and runs them in production: I moved a 7-person studio team onto agent-driven development, automate its business processes with Hermes agents and MCP integrations, and operate a medical-data RAG assistant for client users. I co-built the ticket-driven Claude Code orchestrator that runs my own delivery and self-host open models for my agents. 10+ years of full-stack engineering (Python, TypeScript), 5 of them leading teams.

AI Engineering Projects

LLM-driven development pipeline for a production studio (Talking Birds & Flying Fish, Aug 2025 — Present)

- Moved the studio team onto agent-driven delivery: added the first agent configuration in Aug 2025, and a year later it was in 22 of 33 repositories, with every 2026 project starting from it. Introduced a plan → isolated worktree → independent agent review → human acceptance → ship workflow that the team adopted, and measured the change with an evidence-tagged atlas of git, tracker and PR history.
- Ran studio work through the yokemate orchestrator: 169 agent-written plans across 11 projects, 23 studio tickets merged through the pipeline in its first month.

Hermes agents for business-process automation (Talking Birds & Flying Fish, Apr 2026 — Present)

- Deployed Hermes agents for staff with a shared skill bank I wrote — CG budget estimates from a brief to a priced spreadsheet and client PDF, estimate-template expertise, a freelancer database with privacy rules — connected to Microsoft 365 over MCP and covered by a promptfoo eval suite with LLM-as-judge.
- Built the data layer behind them: a corporate data core designed from an audit of every department's spreadsheets, with its 22-migration schema rolled out by an agent through the Supabase MCP server; an hourly Python sync to SharePoint via Microsoft Graph, in production since Sep 2026; the HR vacation tracker moved onto the database in two weeks (10 merged PRs, tests from 21 to 135).
- Built Microsoft Teams bots on the data core — a company-wide HR bot (leave, worklogs, calendar, policy search) and data bots for management — and a company RAG over SharePoint documents, now in integration.

Aleen — medical-data AI assistant (YokeLoop client, Jun 2026 — Present)

- Run a Hermes agent over an MCP-based RAG system on family health records, with per-user data isolation; 6 active users. Administer the deployment on the client's server and tune the harness and agents.

yokemate — ticket-driven multi-repo orchestrator on Claude Code (YokeLoop, co-built with [Ivan Hilkov](#), Aug 2026 — Present; private, demo on request)

- Runs my delivery across 4 organizations and about 18 repositories: 88 tickets in the first month, 72 merged at a median of 16 hours from plan to merge, 1 in 5 sent back for rework at acceptance review.
- Read-only scout and planner subagents, implementation in an isolated worktree, an independent reviewer subagent and a human merge gate; a command guard lets task sessions run unattended. Grew out of the co-authored [yoke](#) plugin; my rebuild on the pi runtime, [yokemate-pi](#), showed the workflow-first design was the wrong foundation and led to mypi, a memory-first successor (early stage).

Realtime voice agent — 3 iterations (personal R&D, Apr 2026 — Present) · [v1](#) · [v2](#) · [v3](#)

- Cut NVIDIA Audio2Face-3D GPU memory from 8.8 GB to ~0.4 GB by rebuilding its TensorRT engine for single-stream use, and per-utterance facial animation from ~3 s to 81 ms with a persistent C++ inference helper — so STT, TTS, an LLM and avatar animation fit on one 12 GB GPU.
- Evolved from a 3-node hybrid with a Live2D avatar (22 ADRs) to a local-first single host with JSON-Schema tool calls in a Docker sandbox, then to v3 with LangSmith tracing and a single `delegate(agent, task)` tool to coding agents; chose TTS by gates I defined (latency, real-time factor under load, VRAM headroom, Russian quality).

Personal agent stack (personal, Feb 2026 — Present)

- Run Priney, my always-on assistant on Hermes, self-hosted on a Raspberry Pi 5, covering tasks, calendar, email, finances, health, knowledge and travel. Designed its Second Brain — an Obsidian vault of 226 notes with routing rules for where the agent reads and writes, synced through CouchDB and git — and an on-device RAG over it (ONNX embeddings, SQLite vector store, incremental re-indexing).
- Wrote its integrations with self-hosted services: MCP servers for tasks (`sp-cli`) and, earlier, habit and time tracking, alongside finance and bookmark services; plus cron automations — a watchlist monitor with priorities and quiet hours, digests, nightly backups.


- Self-hosted Qwen3.6-35B-A3B as my agents' primary model on a single 12 GB GPU (MoE expert offload, quantized KV cache, 262K context, tool-calling template) and tuned the agent around it — enforced tool use, compression that keeps tool-call history; tested smaller models (LFM, Gemma, Qwen, community fine-tunes) on the GPU and the Pi; built llama-tray to switch model presets.

sp-cli — CLI and MCP server for Super Productivity (Sep 2026)

- 97 commands with an MCP server generated from the CLI itself; built in 2 days by an autonomous agent session on the yoke workflow and verified against a live phone sync.

Key Metrics (July 2025 — September 2026)

- Throughput with a quality denominator: 3,900+ commits and 368 pull requests (341 merged, 93%) across 59 repositories in 15 months; only 12 commits were reverts.
- Context engineering: about a third of commits change only Markdown or agent configuration (plans, journals, skills, CLAUDE.md); about 25 skills and 25 subagent definitions authored and versioned in git.



Drawn from local git history across my personal and work GitHub accounts; GitHub's own graph on this profile misses about 1,350 studio commits made under the work account.

Professional Experience

YokeLoop — Co-founder (product & client delivery) — Apr 2026 — Present

Two-founder company building LLM-agent products for small and mid-size businesses; I own product, specifications and client delivery, my co-founder owns core engineering

- Specified hermes-orchestration, a platform for hosting isolated AI agents on bare metal (a sandbox per user, credentials kept on the host); deploy and operate it for clients, including Aleen and Talking Birds.
- Co-authored the yoke Claude Code plugin and the yokemate orchestrator.

Talking Birds & Flying Fish — Tech Lead → Team Lead → Head of Development — Jul 2024 — Present

Barcelona production studio: interactive installations for corporate conference stands and museum exhibitions; clients include Microsoft. Tech Lead (Jul 2024), Team Lead (May 2025), Head of Development (Jan 2026)

- Grew the web team from 2 to 5 developers and lead them alongside 2 Unreal Engine developers: 40 installations (25 visitor-facing) from 33 repositories — Microsoft conference stands touring 12 cities, game stations, a permanent museum exhibition — with 2,299 tracked tasks and 670 PRs; 1,128 own commits.
- Lead the studio's AI adoption: the agent-driven development pipeline and Hermes business automation above.
- As developer on the museum exhibit, replaced a legacy .NET drawing recognizer with a Python OpenCV service without touching installed hardware; built a 40-case test harness, disproved my own planned fix with diagnostics and moved the remaining error class into the printed form's marker codes.
- Co-developed the shared @tb-ff/web-toolkit (used in 18 of 33 projects) and set up centralized logging (Loki, Grafana).

Sminex — Senior Full-Stack Developer, Team Lead (Business Process Automation) — May 2022 — Jul 2024

- Led the in-house business-process automation team of 3 at a real-estate developer (process, architecture, hiring); shipped 5 web services automating construction and operational workflows, integrated with internal systems and supported in production. TypeScript, React, NestJS, PostgreSQL, Django, Autodesk Forge.

SU-10 — Tech Lead, Construction Process Automation — Aug 2021 — May 2022

- Led a team of 3 automating construction business processes on the Autodesk Forge (BIM) API, from development process to MVP in 10 months; 1st place at the Autodesk Forge Hackathon 2021 (Cloud Collaboration).

RoadAR — Frontend Developer — Feb 2022 — Sep 2022 (part-time)

- Web visualization of point clouds and spatial data from SLAM and ML pipelines (React, deck.gl, WebGL, Cesium).

DataMap Geodata Laboratory — Frontend Web Developer — Mar 2018 — Aug 2021

- GIS web services and geodata collection: 5+ client projects through Upwork, 5 data-collection projects; trained 20+ students and co-organized 2 hackathons.

Earlier roles: AT Consulting — Frontend Developer (2017 — 2018); freelance web development (2015 — 2016); PHP internship (2014).

Technical Skills

- **Agents & LLM:** Claude Code (skills, subagents, hooks), Claude Agent SDK, pi, Hermes Agent, MCP servers, context engineering, orchestrator-worker pipelines with human merge gates, structured outputs and tool calling, LiteLLM
- **Evals, RAG & inference:** promptfoo with LLM-as-judge, test harnesses, RAG (chunking, multilingual embeddings, vector stores, SharePoint via Microsoft Graph), LangSmith, llama.cpp (quantization, MoE offload, speculative decoding), TensorRT, Whisper and TTS models, ComfyUI
- **Stack & workstation:** Python (FastAPI, Django, pytest), TypeScript, React, Electron, Node.js, PostgreSQL, Supabase, SQLite, OpenCV, Docker, GitHub Actions; Arch Linux (Omarchy), Neovim, herdr, pi, Claude Code
- **Languages:** Russian (native), English (fluent, professional working), Serbian (basic)

Education

Information Systems and Technologies — Moscow State University of Geodesy and Cartography (MIIGAiK), 2013 — 2017

Availability

- Open to full-time AI Engineer, Agent / LLM Platform Engineer, Applied AI and Forward Deployed Engineer roles; remote, worldwide; B2B contract or employer of record.